

Annotating single-cell data on a laptop: a 13-method benchmark and practical low-memory workflows, with actinn-jax

Ian Driver*

Abstract

Reference-based cell-type annotation is a routine, high-volume step in single-cell analysis, and most labs run it on a laptop, not a GPU cluster. We present **actinn-jax**, a from-scratch JAX reimplementation of ACTINN (a 4-layer MLP classifier) with a sparse-aware pipeline and a train-once/map-many cached reference model, and use it as the occasion for a neutral benchmark of **thirteen** methods spanning every major class — classical (SVM, kNN), regularized-linear (CellTypist), a carefully-tuned linear pipeline, parameter-free projection (scTOP), marker/correlation (SingleR, scmap), deep probabilistic (scANVI, scArches), an interpretable prototype VAE (ProtoCloud), and a foundation model (scPRINT) — across **six** datasets (within-dataset, cross-dataset and cross-study; 8 to 86 cell types) plus a five-point **scaling sweep to 49k reference cells**, measuring accuracy, macro-F1, ontology-aware concordance, training and inference time, and peak memory on commodity Apple-Silicon hardware. **Accuracy differences among the leading methods are small; their cost differences are not.** The top of the accuracy table is a four-way cluster spanning 0.008 (0.831–0.839), led by a carefully preprocessed **linear pipeline** (normalize → ANOVA → PCA → logistic regression, 0.839) rather than by a deep model, with scANVI, scArches and actinn-jax (0.831) alongside it. **Those same four methods differ by ~130× in inference time** (0.34 s to 89 s) **and 2.4× in peak memory** (1.8 to 4.3 GB). The linear pipeline is the most accurate-per-second option on laptop-sized data, fitting 7× faster than actinn-jax (3.0 vs. 21.9 s). At atlas scale the ordering changes: **ProtoCloud** rises from worst at 3k reference cells (0.722) to best at 49k (**0.976**, vs. 0.936 for actinn-jax and 0.939 for the linear pipeline). No method leads everywhere, and methods separated by less than a percentage point of accuracy are separated by two orders of magnitude in cost. Because several methods now annotate quickly, accurately, and with a usable handle on unknown cells, the useful contribution is not another ranking but an account of **which tool fits which job, and what the surrounding workflow looks like**. Every method is run on every dataset through one harness on identical splits, the panel deliberately includes the baselines most likely to beat a small MLP, and the external validation is somebody else’s benchmark — **Open Problems (OP) label_projection**, whose datasets, metrics and ranking we did not choose. We report the results as a decision aid (accuracy-per-second, accuracy-per-byte, and behaviour from 3k to 49k reference cells), then use actinn-jax to demonstrate **end-to-end workflows that a low, flat cost profile makes practical**, on human data and — via a hierarchy derived from the Cell Ontology, which removes the GPU step from building a reference — on mouse: annotating an unknown human dataset from a **shipped census-scale reference** (~800 cell types, calibrated abstain) with no training; **tissue-aware refinement** that halves spurious cross-tissue labels on a liver query; a **broad→refined hand-off** to a small focused reference (held-out cross-study liver 0.23/0.58 → 0.72/0.86, exact-match / Cell-Ontology-aware); within-cell-type resolution (hepatocyte zonation); and a **cluster-level novel-cell-type screen** that recovers a withheld pulmonary ionocyte population and its marker ASCL3. What enables this is inference that is **sub-second and flat** — independent of reference size and cardinality — over a cached reference that is trained once and reused, at the lowest peak memory of the methods we carried to atlas scale (6.1 vs. 13.2 GB for the linear pipeline at 49k cells). The components are not unique: ProtoCloud provides uncertainty, gene attribution and

a retraining-based refinement path, and **Pan-human Azimuth** ships a pretrained, calibrated, hierarchical pan-human annotator that runs on a laptop — a better-resourced broad pass than ours, and one we can **distill into a broad-pass model of comparable accuracy at $\sim 9\times$ its throughput, built with no GPU and no labeled data**. What the workflow adds is the **hand-off into a label set the broad model was never trained on**: refinement against a user’s own focused reference, and resolution below any fixed typology’s leaves.

Independent Researcher, Detroit, MI, USA

* Correspondence: driver.ian@gmail.com

1 Introduction

Annotating cell types by mapping to a labeled reference is one of the most frequently run operations in single-cell RNA-seq. The method landscape is large (survey), but published comparisons (Abdelaal 2019, Fu 2024) emphasize accuracy and rarely report the axis that decides what a working scientist actually runs on their own machine: **wall-clock time and memory on commodity hardware**, without a GPU. Foundation models (scGPT, Geneformer, scPRINT) push accuracy in some settings but need GPUs and minutes-to- hours, and — as we and others find — their *zero-shot* label predictions underperform a small model trained on curated reference labels.

We ask a practical question: **for a given annotation job on commodity hardware, which method should you actually run, and what does the surrounding workflow look like?** Several methods now annotate quickly and accurately and give a usable signal on unknown cells, so the shortage is not another leaderboard but guidance on fit-for-purpose — accuracy per second, accuracy per byte, and how each behaves as the reference grows. We answer with a neutral 13-method benchmark plus a set of demonstrated workflows, held to a single standard: every method runs in its own environment through the same harness, on the same splits, scored by the same metrics, and reported as measured — including where actinn-jax is not the leader.

How the comparison was constrained. A benchmark written by a method’s own author has a well-known failure mode: the datasets, baselines and metrics that flatter the method are the ones that get reported. We tried to design that freedom out rather than promise restraint. Every method runs on **every** dataset, so none appears on a favourable subset, and the two methods outside that matrix are named with reasons (§2.2). The panel was chosen to include the baselines most likely to beat a small MLP — a carefully tuned linear pipeline, scTOP, ProtoCloud — and each of them does beat it somewhere; our own classical tier is left untuned while the linear baseline is tuned, which is the unfavourable direction (§5). Metrics are fixed across all methods and reported in full. And the **external validation is somebody else’s**: Open Problems `label_projection` (§3.7) sets the datasets, the metrics and the ranking, and we did not choose any of them. Where results contradict our own earlier claims or our method’s interests — ProtoCloud overtaking actinn-jax at atlas scale, a memory advantage that stays a fixed factor instead of growing with the data, a better-resourced broad annotator than our own — they are reported in those terms.

On adding another method to a crowded field. We are adding one, and we think it earns its place: the original ACTINN no longer installs or runs on current toolchains, and the reimplementation’s flat, sub-second, memory-bounded inference over a cached reference is what makes the multi-stage workflow of §3.4 practical on a laptop. What we are *not* doing is the move that makes method papers hard to trust — concluding that the new method wins. It does not: on accuracy it sits inside a four-way cluster it does not lead, the panel was assembled to include the baselines most likely to beat it, and they do. A field already full of working methods (xkcd 927) is a good

reason to be specific about what a new entry is *for* — here, a cost profile that lets several models be chained — rather than to claim it supersedes what exists. Where a better-resourced model already exists, the productive move is to build on it (§3.4).

Contributions.

1. A modern, dependency-light (no TensorFlow) JAX reimplement of ACTINN with sparse preprocessing, chunked atlas-scale prediction, and a cached reference model.
2. A neutral **13-method** $\times$ **6-dataset** benchmark of accuracy, speed, and memory on Apple Silicon, with a rejection/abstain analysis and a **five-point scaling sweep to 49k reference cells**. The panel deliberately includes the baselines most likely to beat a small MLP — a tuned linear pipeline, scTOP, and ProtoCloud — and the results are reported as measured (§3.1, §5). The sweep is also what sets actinn-jax’s minibatch policy: at ACTINN’s fixed batch of 128, a 47k-cell reference costs $\sim 24k$ tiny update steps, so actinn-jax scales the batch with the reference above $\sim 12.8k$ cells and keeps 128 below it. Measured fit and predict curves are in §3.3.
3. **Demonstrated end-to-end workflows** for jobs where cost is the binding constraint, all CPU: annotate an unknown human dataset from a **shipped census-wide ~ 800 -type reference** with a calibrated abstain and no training; **tissue-aware refinement** (halves spurious cross-tissue labels on a liver query); a **broad $\rightarrow$ refined hand-off** to a small focused reference (cross-study liver $0.23/0.58 \rightarrow 0.72/0.86$); within-cell-type resolution (hepatocyte zonation); and a **one-call novel-cell-type screen** (recovers a withheld pulmonary ionocyte population and its marker ASCL3). We are precise about what is not ours: ProtoCloud offers uncertainty, attribution and a retraining-based refine, and Pan-human Azimuth ships a pretrained hierarchical pan-human annotator with trained abstention, so the broad pass itself is not a distinguishing feature — we distill that model into a broad-pass reference rather than compete with it. What is distinct is **refinement into a label set no pretrained model carries** — the user’s own focused reference, and states below a fixed typology’s leaves. The same route builds a **pan-mouse reference** (453 types / 85 tissues; 0.638 ontology concordance on two withheld datasets) in 17 s of CPU, using Cell Ontology lineage in place of a foundation-model embedding — which beats the embedding-derived hierarchy on the human corpus it was validated against, and needs no GPU at any stage.
4. An **independent external validation** on Open Problems `label_projection`, with a controlled same-hardware speed/memory comparison, and a set of cheap ablations that characterize *when* the gene-space model gains or loses to heavier methods (input standardization; a reference-guided gene budget; a foundation-model negative control).

2 Methods

2.1 actinn-jax

actinn-jax reimplements ACTINN (Ma & Pellegrini 2020) — a 4-layer fully-connected network (100/50/25 hidden units, ReLU, softmax) trained with Adam — in JAX/optax (Bradbury 2018), replacing the original’s TensorFlow-1.x graph/session code. Key engineering:

- **Sparse-aware preprocessing:** the count matrix is never densified in full. CP10k+log2 normalization and the expression / coefficient-of-variation gene filter run on the sparse matrix; only the selected-gene columns of each minibatch are made dense.
- **Cached reference model:** train once, write the fitted model to disk, reload it and map

many queries — the amortized cost of repeated annotation against a fixed reference drops to the inference cost alone.

- **Chunked prediction** for atlas-scale queries, and raw counts are picked up automatically from the slot CELLxGENE files conventionally store them in.
- **Optional reference-fit standardization:** z-score each selected gene by the reference’s frozen mean and standard deviation and apply that to the query — a cheap domain alignment (§3.7). Off by default; when on, the scaler is saved alongside the model and travels with it, and models saved without it still load.

The reimplementaion reproduces the original’s accuracy to within repeat noise while being **substantially faster and lighter and, unlike the TensorFlow-1.x original, installable on current Python/ML environments** — the original’s graph/session code no longer runs on a modern stack without pinning a years-old TensorFlow. Because the two are the same model at equal accuracy, we do not carry the original as a separate benchmark row; the engineering win is the point, not an accuracy comparison.

2.2 Benchmarked methods

method	source	tier	engine	rejection	runtime
actinn-jax	Ma & Pellegrini 2020	classical	JAX MLP	confidence threshold	JAX
SVM	Pedregosa 2011	classical	linear SVM (SGD)	—	scikit-learn
kNN	Pedregosa 2011	classical	k-nearest neighbors	—	scikit-learn
CellTypist	Domínguez Conde 2022	linear	L2 logistic regression	prob threshold	scikit-learn
linear-anova-pca	Pedregosa 2011 †	linear	normalize → ANOVA → PCA(220) → logreg	prob	scikit-learn
scTOP	Yampolskaya 2023	parameter-free	rank z-score class-average projection	—	NumPy
SingleR	Aran 2019	correlation	Spearman + fine-tuning	—	R
scmap-cluster	Kiselev 2018	correlation	centroid cosine	yes (unsigned)	R
scANVI	Xu 2021	deep	scVI semi-supervised VAE	prob	scVI
scArches	Lotfollahi 2022	deep	scANVI reference surgery	prob	scVI
ProtoCloud	Guo & Ding 2026	deep	prototype VAE + LRP attribution	ambiguity flag	PyTorch
scPRINT	Kalfon 2025	foundation	pretrained transformer, zero-shot	—	PyTorch

method	source	tier	engine	rejection	runtime
Pan-human Azimuth	Sarkar et al. 2026	pretrained	8-level hierarchical NN, fixed 382-leaf typology	trained Unassigned	TensorFlow

Table 1. The benchmarked methods: model family (*tier*), the engine each one actually runs, whether it can decline to call a cell (*rejection*), and the runtime stack it was run on; each has its own pinned environment (§2.5). Bold marks the methods added to this panel here.

† `linear-anova-pca` has no method paper: it is a baseline assembled here from scikit-learn components, tuned deliberately to be the strongest simple competitor we could build (§1).

Eleven of the thirteen are trained on each dataset’s own reference and run the full accuracy/cost matrix of §3.1 on all six datasets, so none is compared on a favourable subset. The remaining two are **pretrained annotators with fixed label vocabularies**, which changes what can be measured rather than excusing them from measurement:

- **scPRINT** is a zero-shot foundation model over a fixed **Cell Ontology (CL)** vocabulary — CL being the standard controlled vocabulary of cell types, used throughout as the common ground between datasets that name the same cell differently; it cannot emit labels outside that set and skips symbol-keyed datasets.
- **Pan-human Azimuth** ([Sarkar et al. 2026](#)) is a supervised hierarchical classifier over a single organism-wide typology — 8 levels, 382 leaves, ~7M parameters, a fixed 5,055-gene panel, trained on 9.7M curated cells, with abstention *learned rather than thresholded* — the model is trained to answer **Unassigned**, instead of having a probability cutoff applied to its output afterwards. It is the closest published counterpart to this paper’s broad pass, and the teacher we distill in §3.4.

Neither predicts into the dataset’s label strings, so **exact-match accuracy against those strings is a vocabulary artifact, not an accuracy signal** — regulatory T cell → Treg cell and non-classical monocyte → CD16 monocyte are correct answers that score zero. Both are therefore scored **ontology-only** (§2.4) and reported in their own block in §3.1, on the three datasets where reference and query both carry Cell Ontology ids. They appear again in §3.6 (foundation-model zero-shot) and §3.4 (Azimuth as broad pass, and as distillation teacher).

scPred is excluded outright: unmaintained; it requires harmony’s removed `HarmonyMatrix` API and older harmony versions do not compile on the current toolchain.

2.3 Datasets

dataset	source	split	tissue	#types	genes	notes
lung_intra	Travaglini 2020	within-dataset	lung (Krasnow)	46	Ensembl	300 cells/type ref
lung_cross	Sikkema 2023 → Travaglini 2020	cross-dataset	lung (HLCA → Krasnow)	46	Ensembl	different lab/protocol

dataset	source	split	tissue	#types	genes	notes
liver_intra	Edgar 2026	within-dataset	liver (HLiCA)	36	Ensembl	150 cells/type
liver_cross	Edgar 2026	cross-study	liver (HLiCA)	34	Ensembl	train 6 studies → test withheld study
blood_gut_intra	Alegbe 2026	within-dataset	blood + gut (IBDverse)	86	Ensembl	high cardinality; no CL ids
pbmc	10x Genomics 2016	within-dataset	PBMC (pbmc3k)	8	symbols	small-n; scPRINT skips (symbols)

Table 2. The six benchmark datasets. *split* separates within-dataset holdouts from cross-dataset and cross-study transfer; *genes* records whether the matrix is keyed by Ensembl ids or by gene symbols, which decides who can read it.

The blood+gut set is a subsample of **IBDverse** (Wellcome Sanger Institute; blood, terminal ileum and rectum from 421 individuals), included here because 86 fine-grained labels make it the high-cardinality stress case for the panel. It carries no Cell Ontology ids, so it contributes to accuracy and macro-F1 but to no ontology-scored result.

The set spans the three generalization regimes (within-dataset, across datasets, across studies), an order of magnitude in cell-type count (8→86), and both gene-ID conventions.

2.4 Metrics

Per (dataset, method, repeat): **accuracy**; **macro-F1**, unweighted over classes so rare types count as much as common ones; **ontology-aware concordance**; **fit time**, **predict time**, and **peak RSS** (resident set size — the maximum physical memory the process held).

Ontology-aware concordance credits a call that is the same node, an ancestor, or a descendant of the truth in the Cell Ontology. It exists because cell-type vocabularies disagree about granularity more than about biology: **periportal region hepatocyte** against **Hepatocyte** is a right-lineage-wrong-depth call, which exact match scores as a total failure and a working scientist would accept. It is the **only** metric that compares methods with different label vocabularies (§2.2), and it is reported only where both reference and query carry CL ids — so not on **blood_gut_intra**, and not on the symbol-keyed **pbmc**.

Concordance depends on the ontology release, since ancestor sets change between them. All numbers here use the **Cell Ontology basic release of 2026-06-08** (sha256 73996c63...), resolved once and reused by every scoring pass. This is not a hypothetical concern: re-running **lung_cross** against a later-fetched release reproduced the splits exactly (14,390 reference / 13,550 query cells) and still moved concordance by 0.003, with no model changing. That release, its checksum, and pinned versions for all six benchmark environments are recorded as lockfiles in the [benchmark repository](#), which also carries a command that reports any drift from them.

Exact-match accuracy is reported for the ten reference-trained methods and, on `lung_cross`, flagged: reference and query use different label vocabularies there, so its exact accuracy is a vocabulary artifact and the ranking, not the level, is what transfers.

2.5 Harness, isolation, and hardware

Each method runs in **its own virtual environment, as a separate process**, because the dependency sets are mutually unsatisfiable — scVI-based methods, TensorFlow/Keras (Pan-human Azimuth), R (SingleR, scmap) and the JAX core cannot coexist in one interpreter. One driver builds the reference/query split **once per dataset** from a fixed seed and hands the identical pair to every method, so no method sees a different split. **repeats = 3**; scPRINT and Pan-human Azimuth run once each, being deterministic given a fixed query and dominated by a single forward pass.

Resource accounting is per subprocess: fit and predict are timed separately and peak RSS is sampled by a monitor thread in the child, so one method’s footprint cannot be attributed to another.

Hardware: Apple Silicon (M-series), CPU for the classical, linear, correlation and pretrained tiers and for actinn-jax; Apple MPS for the deep and foundation tiers. No discrete GPU, no cloud — the regime the benchmark is about. Math-library threads are left uncapped, so wall-clock is what a practitioner sees; an optional thread cap exists for stricter reproducibility. §3.7 repeats the comparison on one cloud CPU instance, where matched concurrency matters and is discussed (§2.10).

2.6 Building a broad reference

The shipped broad-pass references (§3.4) are built once, offline, by a three-stage pipeline driven by one command (in the [benchmark repository](#)).

Sampling. Every primary cell for one organism in the CELLxGENE Census [CZI Census 2025](#) — primary meaning the copy the census designates canonical, so cells deposited in more than one study are counted once — stratified by cell type at a fixed cap per type: 40 for `broad_human_v1`, 60 for the later human and mouse pulls. Types with fewer than 12 cells are dropped, as are labels that carry no information (“unknown”, “native cell”, “eukaryotic cell”, “animal cell”). The pull is checkpointed every ten datasets and retried, since transient object-store read errors are routine at this scale. Genes are carried under both Ensembl ids and symbols, because a symbol-keyed consumer cannot use an Ensembl-only pull. The census “stable” alias is a moving pointer, so the release it resolved to is recorded; every reference here comes from **census 2025-11-08**.

Coarse hierarchy — two interchangeable routes. The model is a coarse classifier over groups plus one fine classifier per group, and the grouping comes from either:

1. **Foundation-model embeddings.** scPRINT (medium-v1.5) embeds the reference once in 4,000-cell chunks; per-cell-type centroids are Ward-clustered into roughly $\sqrt{(\text{number of types})}$ groups, with a floor of eight. scPRINT’s QC drops cells, so the embedding is not positionally aligned to the input and the label is carried through the model with the vectors. This is the only step that wants a GPU/MPS.
2. **Cell Ontology lineage.** Each cell type is described by a binary indicator over the CL terms it descends from; terms appearing in fewer than two types, or in all of them, are dropped as uninformative; the same Ward clustering is applied to those vectors. Free, deterministic, and species-independent — which is what makes the pan-mouse reference possible, since

the distillation teacher is human-only and scPRINT’s mouse support is untested here. §3.4 compares the two routes against a random-grouping control.

Training and calibration. A 4,000-gene panel of **highly variable genes (HVGs)** — the genes that differ most across cells, and so carry most of the signal for telling types apart — is selected on the reference; the coarse and per-group fine models are trained on it. Abstain calibration holds out **10% of cell types entirely** as out-of-distribution (OOD — types the model was never trained on, standing in for the novel populations a real query contains), plus a 20% within-type test split, and sweeps the confidence threshold to trade coverage against accuracy on kept cells — the table each shipped reference reports.

Shipping and verification. Each reference ships with a build record: census release, corpus sizes, hierarchy route, organism and calibration table. Because a rebuild can train cleanly and annotate *worse*, the previous reference is backed up and the new one is scored on a held-out atlas before it is kept. For `broad_mouse_v1` the test set is carved from the census itself: two datasets are excluded from the reference sample and re-pulled as the query, so “held out” means held-out *datasets* and not merely held-out cells.

2.7 Distilling a pretrained annotator

A broad reference can also be built without labels and without a GPU, by taking both the label vocabulary and the hierarchy from an existing pretrained annotator. Keras/TensorFlow and JAX cannot share a process, so this runs in two stages.

Teacher pass. Pan-human Azimuth labels a corpus of raw counts through its low-level API in batches of 8,192, which loads the weights once rather than once per call. All eight levels are retained. Where the package cannot reconcile a refined level with its coarser ones it blanks that column; those cells still carry a deepest-level call, so we fall back to it rather than discard them — dropping them would distill only the cells the teacher found easy. Predictions are mapped to CL ids through the package’s own crosswalk.

Student pass. The student is a hierarchical `actinn-jax` reference whose **classes** are the teacher’s refined fine labels (including its trained **Unassigned** quality-control class) and whose **hierarchy** is the teacher’s own broad level — so the coarse→fine structure is inherited rather than rediscovered. 4,000-gene HVG panel; classes with fewer than 8 cells dropped, since a class that cannot be split into train and test makes fidelity depend on which side its cells landed.

Corpus. Three local atlases (lung, liver, blood+gut) capped at 600 cells per label, plus a census-wide pull (≤ 60 per type). The atlases’ own labels are used **only for scoring**, never for training.

Evaluation. Two arms answer different questions. *In-corpus* holds out a stratified 25% of cells and measures **fidelity** — how often the student reproduces the teacher, exactly and ontology-equivalently. *Held-out atlas* withholds an entire atlas and measures whether the student generalizes to data the corpus never covered. Both arms also score student and teacher against the atlases’ own labels, which separates distillation loss from teacher error. Note that once the census pull is in the corpus, a withheld *atlas* is no longer a withheld *tissue*.

2.8 Refinement, abstention, and novelty

Four mechanisms operate on a trained reference at inference time, all label-free.

- **Confidence abstention.** Cells whose final probability is below a threshold are relabeled “unknown” rather than forced to the nearest class. Calibrated per reference (§2.6).
- **Refinement to the query.** The reference’s class set is masked to those classes the query’s own predictions actually support, and the softmax renormalized — no ground truth, no retraining.
- **Refinement to a tissue.** Classes are masked to those the census records in the given tissue, while **pan-tissue** types (immune, endothelial, stromal) stay available everywhere, so liver-resident T cells are unaffected. The per-class tissue map is baked into the shipped reference; the tissue is given explicitly or read from the query’s own metadata, with common synonyms (PBMC→blood, hepatic→liver) recognized. A tissue outside the reference’s vocabulary imposes no filter rather than an empty one.
- **Cluster-level novelty screen.** A cell type absent from the reference appears not as scattered uncertain cells but as a *coherent group the reference cannot confidently explain*. Clusters that are both large enough to be a real population and predominantly low-confidence are flagged, with marker genes reported for each. This is distinct from per-cell abstention.

Within-type state. The same machinery resolves structure below a cell-type label: a hepatocyte zonation model trained on portal→central zone labels, scored by within-one-zone agreement across donors and datasets.

2.9 Scaling protocol

The reference is subsampled to five sizes up to **49k cells** with the query held fixed, so that accuracy, fit time, predict time and peak memory are functions of reference size alone. Reported per method. The sweep exists to find where one method overtakes another as the reference grows, which no single reference size can show.

2.10 External validation protocol

Open Problems label_projection [Open Problems 2024](#) supplies the datasets, metrics and ranking — none chosen by us. `actinn-jax` is packaged as a `viash 0.9.7` component declaring the normalization it expects, and run through the project’s own Nextflow workflow. We additionally built components for four baselines (linear-anova-pca, scTOP, SVM, CellTypist) so the same-hardware comparison covers more than the leaderboard’s own set.

Two caveats govern how those numbers may be read. The leaderboard aggregates runs from **heterogeneous cloud hardware**, so its timings are not a controlled cost comparison — hence the same-hardware run: every method on one AWS CPU instance. And in that run, **wall-clock is only meaningful at matched concurrency**: with several tasks per box, `%cpu` spanned 49% to 2338% across methods, so cross-method runtime ratios measured at high concurrency rank threading and scheduler contention as much as algorithm. Cost is therefore reported only from the matched-concurrency run (Table 10); the four components we added to the harness are compared on accuracy alone (Table 11), which does not depend on the machine.

2.11 Ablations

Three cheap interventions characterize *when* a gene-space MLP gains or loses against heavier methods, all selectable without test labels:

- **Input standardization** — z-score each selected gene by the reference’s frozen mean/std

and apply to the query (a scArches-style domain alignment). Opt-in, because it shifts the probability calibration the abstain thresholds are tuned against.

- **Gene budget** — Open Problems feeds every method 1,000 HVGs; we sweep wider panels and select per dataset by **held-out reference cross-validation**, which is label-free with respect to the test set.
- **Protein-embedding featurization** (negative control) — a CPU-only featurization in the style of **Universal Cell Embeddings (UCE)** [Rosen 2026](#) — an expression-weighted mean of ESM2 [Lin 2023](#) protein-language-model gene embeddings — to test whether a foundation model’s value survives a cheap pooling shortcut.

3 Results

3.1 Accuracy and cost

Accuracy is the mean over the **five shared-vocabulary datasets** (lung_cross excluded — its exact accuracy is a vocabulary artifact, see the † note below; means over all six are ~0.08 lower for every method, with the identical ranking); macro-F1 and ontology are means over all six; cost is the mean per query (Table 3):

method	acc (5)	macro-F1 (6)	ontology (6)	fit (s)	predict (s)	peak MB
linear-anova-pca	0.839	0.699	0.808	3.0	0.34	4294
scANVI	0.833	0.697	0.809	0.0*	89.2	2099
scArches	0.832	0.699	0.805	54.7	21.4	1819
actinn-jax	0.831	0.683	0.811	21.9	0.67	2399
CellTypist	0.823	0.691	0.805	60.7	0.66	1569
SVM	0.810	0.680	0.797	66.5	0.07	1415
ProtoCloud	0.790	0.649	0.778	176.3	2.14	1893
SingleR	0.770	0.652	0.750	0.3	62.1	3541
kNN	0.770	0.622	0.769	0.4	0.15	1483
scTOP	0.739	0.619	0.703	1.1	1.07	1928
scmap-cluster	0.646	0.550	0.771	0.4	13.1	8073

Table 3. Accuracy and cost together, ordered by accuracy. Accuracy is the mean over the five shared-vocabulary datasets, macro-F1 and ontology concordance are means over all six, and cost is the mean per query. Bold marks the best value in a column.

*In Table 3, scANVI does most of its work in one train+predict pass, attributed to predict.

The top four methods span **0.008 in accuracy**, ~**130× in predict time** (0.34 s to 89 s) and 2.4× in memory. scANVI sits within 0.002 of second place and **134× slower than actinn-jax** at tied accuracy. The linear pipeline is both the most accurate and the fastest to fit, and pays for it in memory — 4294 vs 2399 MB — because ANOVA/PCA densify a cells × genes matrix while actinn-jax stays sparse; that ~1.8× ratio widens to ~2× at atlas scale (6.1 vs 13.2 GB at 49k cells, §3.3). The two profiles suit different jobs: one-shot labelling, versus a reference **trained once and reused**, where a cached model pays fit once while the linear pipeline refits scaler/PCA/classifier for every query. ProtoCloud’s 146 s CPU fit buys nothing at this reference size — and buys the top accuracy at atlas scale (§3.3).

Pretrained annotators, scored ontology-only. The two methods with fixed label vocabularies (§2.2, §2.4) cannot appear in Table 3, since they do not predict the dataset’s label strings. They are scored by ontology concordance on the three datasets where reference and query both carry CL ids, with actinn-jax on the same splits for reference:

	lung_intra	liver_intra	liver_cross	predict (s)
actinn-jax (reference-trained)	0.917	0.846	0.731	0.27–0.37
Pan-human Azimuth (pretrained)	0.700	0.521	0.408	2.2–3.4
scPRINT (zero-shot)	0.201	—	—	62.3

Table 4. Pretrained annotators, scored by ontology concordance only, on the three datasets where reference and query both carry Cell Ontology ids — with reference-trained actinn-jax on the same splits for scale. Not a like-for-like comparison; see the text below.

This is not a like-for-like comparison and should not be read as one. actinn-jax is trained on a reference drawn from the same data and scored in its own vocabulary; Pan-human Azimuth has never seen these datasets and answers in a different one. A reference-trained model *should* win. What the block does establish is the distance between a curated supervised annotator and a zero-shot foundation head — 0.700 against 0.218 on lung — and it sets up §3.4, where Azimuth serves as a broad pass in its own right and as a distillation teacher. Peak memory is indistinguishable between actinn-jax and Azimuth (1.65–2.1 GB on every dataset).

The top of the accuracy table is a **four-way tie within 0.008** — the tuned **linear pipeline (0.839)**, scANVI (0.833), scArches (0.832) and actinn-jax (0.831) — led by the linear pipeline rather than by a deep model. CellTypist (0.823) sits just outside it. actinn-jax has the best ontology-aware concordance (0.811), marginally ahead of scANVI’s 0.809 — a gap well inside repeat noise, so read it as a tie rather than a lead. ProtoCloud (0.790) and scTOP (0.739) sit below that cluster on these subsampled references: both need conditions these splits do not provide — ProtoCloud is data-hungry and becomes the strongest method once given a real atlas (§5, SCALING_MEMORY.md), while scTOP is built for small, low-cardinality problems. Per dataset (accuracy):

dataset	actinn-jax	best method (acc)
lung_intra (46 types)	0.894	ProtoCloud 0.932
lung_cross (cross-dataset)†	0.358 exact / 0.752 ontology	ProtoCloud 0.374 / 0.791 ontology
liver_intra (36 types)	0.802	linear-anova-pca 0.804
liver_cross (cross-study)	0.686 exact / 0.731 ontology	actinn-jax
blood_gut (86 types)	0.860	linear-anova-pca 0.902
pbmc (8 types)	0.913	scArches 0.931

Table 5. Per-dataset accuracy: actinn-jax against whichever method leads that dataset. † on

lung_cross the exact-match score is a vocabulary artifact, so ontology concordance is the meaningful column.

No single method leads everywhere: the linear pipeline and ProtoCloud each take two datasets, actinn-jax and scArches one apiece. actinn-jax leads the cross-study liver split — the hardest generalization regime here, and the one closest to real reference mapping — and is second on within-dataset liver, while trailing on lung, on the 86-type blood+gut set (−4.2 pt to the linear pipeline), and on small-n pbmc. The spread across the leading methods on any one dataset is 1–4 points. scmap-cluster is consistently the weakest (0.23 macro-F1 on liver_cross).

† **lung_cross exact accuracy (~0.35 for every method) is a label-vocabulary artifact, not a transfer failure.** HLCA (reference) and Krasnow (query) were annotated independently and share only **20 of 46 cell-type names**; 26 of Krasnow’s types are absent from HLCA’s vocabulary, so no classifier can emit the exact string. The mismatch is granularity: HLCA’s finer taxonomy has *alveolar / elicited / lung macrophage* but **no generic macrophage** (CL:0000235), while Krasnow labels many cells *generic macrophage* — so an HLCA-trained model predicts *lung macrophage* (CL:1001603) for them, exact-wrong but ontology-correct (*lung macrophage is-a macrophage*); likewise Krasnow’s generic *endothelial cell* (CL:0000115) has no HLCA counterpart. **Ontology-aware concordance — which credits a same-lineage call — is ~0.75 for the trained methods (0.67–0.79 across the full panel, highest for ProtoCloud),** so cross-dataset transfer actually works; exact-match just conflates classifier error with vocabulary/granularity mismatch. This is precisely why we report ontology concordance, and it is a general benchmarking pitfall: **cross-dataset exact accuracy between independently-annotated atlases is not a meaningful accuracy signal.** The other five datasets use a single shared label vocabulary for reference and query (a split of one atlas, or HLiCA’s harmonized annotations), so their exact scores are unaffected.

3.2 The accuracy $\times$ speed frontier




Figure 1. Accuracy against total wall time on liver_intra. The tuned linear pipeline holds the fast-frontier point; actinn-jax sits just inside it at this reference size, and what separates it is visible only on the scaling axes of Figure 2.

The **tuned linear pipeline** holds the fast-frontier point: highest matrix accuracy (0.839) at the lowest fit time of any accurate method (3.0 s), placing it above and to the left of actinn-jax, SVM, kNN and CellTypist. The deep methods buy little or nothing on accuracy here — at or below the linear pipeline on every dataset except lung — for 30–130 $\times$ the inference cost. actinn-jax sits just inside the frontier at this reference size; what separates it is not visible on a fixed-size plot but on the scaling axes of §3.3 and §5: predict time that stays flat as the reference grows, and $\sim 2\times$ lower peak memory that holds to atlas scale. The figure is drawn from the liver_intra panel and depicts the ranking of Tables 3 and 5.

3.3 Scaling




Figure 2. Fit and predict time against reference size and label cardinality. Fit time grows for every trained method; predict time stays flat and sub-second across the whole range.

Training time grows with reference size and with #cell types for all trained methods — actinn-jax’s fit goes 2.6 s → 17.4 s → 23.5 s as the reference grows 965 → 14.8k → 24k cells, below CellTypist (6 s → 71 s → 116 s) and SVM (4 s → 51 s → 80 s) at every size in the sweep; kNN is lazy (fit is trivial but it pays at predict and is the least accurate).

The other axis is the one that matters for the workflows: **predict time stays flat and sub-second across the entire range — 0.08–0.33 s whether the reference has 1k or 24k cells, or 5 or 86 types.** Inference cost does not grow with reference size or cardinality, and with the train-once/map-many cache the fit cost is paid once and the flat predict cost is all that recurs — the regime that matters when a reference is reused across many queries. (The scaling script’s memory column is process-cumulative and not a clean per-size measurement; per-method peak memory is reported from the isolated-subprocess matrix in §3.2.)

3.4 The broad→focused annotation workflow

What actinn-jax is built around is not a single classifier but a **two-pass workflow** that matches how annotation is actually done: get a broad call fast, then sharpen it where it matters. Because every stage is a cached `ReferenceModel` with sub-second, memory-bounded inference (§3.1, §3.3), the whole workflow runs on a laptop.

One query, two passes: the broad pass routes, the focused pass resolves




Figure 3. The two passes on the withheld HLiCA liver study (3,396 cells), same UMAP throughout. *Left:* the shipped census reference spreads 137 of its 798 labels over the query with little correspondence to cluster structure — ontology concordance **0.34**. Its job is to establish that this is liver and route accordingly, not to name subtypes. *Middle:* a 36-type liver reference trained on the six non-withheld studies re-annotates the same cells at **0.73**, and its labels now track the clusters. *Right:* the study’s own labels, on the shared palette, for comparison with the middle panel. Both models are leakage-free with respect to this query: the focused reference is trained only on the other six studies (the *shipped* liver reference includes this study and would score 0.94 here, which is why it is not the model plotted). Abstention is switched off in all panels so the numbers measure labelling rather than coverage; §3.5 covers abstain separately.

The broad pass — census-scale. A shipped ~800-type human reference built from the CEL-LxGENE census gives any query a first-pass annotation across the whole body, with a **calibrated abstain threshold** so out-of-reference cells are flagged rather than force-labeled (see the smooth abstain curve of §3.5). This is the “what am I looking at” pass — broad coverage, deliberately not fine-grained.

The broad pass can be built from a model instead of from labels. The census route above needs a foundation model on a GPU to discover its hierarchy. A pretrained pan-human annotator already has one — Pan-human Azimuth publishes an 8-level typology with every node mapped to a Cell Ontology term — so labelling a corpus with it and training actinn-jax on those labels transfers both the vocabulary and the structure. That build needs **no GPU and no labeled input**, only raw human counts: under ten minutes of CPU on 85k cells drawn from three atlases plus a census-wide sample. On 3,396 withheld cross-study liver cells:

broad-pass model	classes	ontology	cells/s
census-built (<code>broad_human_v1</code>)	798	0.338	2,962
Pan-human Azimuth	382	0.380	1,076–1,563
distilled (<code>panhuman_distill_v1</code>)	324	0.406	8,937–10,021

Table 6. Three broad-pass entry points on 3,396 withheld cross-study liver cells: the census-built reference, the Pan-human Azimuth teacher, and the model distilled from that teacher.

This makes the entry point both better and $\sim 3\times$ faster than building it from the census directly, and it answers in a harmonized, ontology-mapped vocabulary rather than raw census strings. Withholding a whole atlas from the corpus, the distilled model tracks the one it was distilled from to within **1.5 points** on lung (0.695 vs 0.710) and **3.0** on liver (0.481 vs 0.511). Distillation itself therefore costs very little: on data neither model was built for, the student stays within a few points of its teacher. What holds the student back is which tissues and studies its training corpus happened to contain — widening the corpus should move these numbers, changing the distillation procedure should not. Two caveats keep this modest. The 0.406/0.380 ordering is one query, and both actinn-jax models draw on a census sample that may include these studies while Pan-human Azimuth does not, so read the distilled model as *comparable to* the one it distills rather than better. And it does **not** inherit that model’s trained abstention: its confidence separates right from wrong poorly (keeping the 90.5% of cells it calls with probability at least 0.5 moves concordance only $0.406 \rightarrow 0.427$), so the calibrated broad-pass abstain of §3.5 belongs to the census-built reference until this one is recalibrated.

A second organism, and a hierarchy that needs no GPU. Neither route above transfers to mouse: Pan-human Azimuth is human-only, so there is no teacher to distill, and the census route wants scPRINT on a GPU with mouse support we have not tested. But the census already labels every cell with a Cell Ontology term, and CL encodes the relation the embedding clustering is recovering — which cell types are kinds of the same thing. Describing each type by the CL terms it descends from and clustering *those* (same Ward linkage, ontology indicator vectors in place of embeddings) gives a hierarchy that is free, deterministic and species-independent.

That substitution needs a control, since any grouping might help. Built from one human corpus (51,346 cells / 867 types) and scored on the held-out lung atlas (per-run numbers in the [benchmark repository](#)):

coarse hierarchy	ontology
Cell Ontology lineage	0.616
flat, no hierarchy	0.547
random grouping, same group sizes	0.539

Table 7. Coarse-hierarchy ablation, ordered by accuracy. Ontology concordance on the held-out lung atlas for a Cell Ontology lineage hierarchy, no hierarchy at all, and a random grouping with the same group sizes as the ontology one.

Random grouping scores no better than no grouping at all (0.539 against 0.547), which rules out the simplest explanation for the gain: that splitting 867 types into 29 smaller problems makes each fine classifier’s job easier, something any split would do. What matters is *which* types share a group. Each group has its own fine classifier and the coarse call decides which one ever sees the cell, so a group pays off only if the coarse classifier can recognize it *and* the right answer is inside it. Lineage groups manage both; groups of the same sizes drawn at random manage neither.

On that basis, **broad_mouse_v1**: 27,026 cells $\rightarrow$ **453 cell types across 85 tissues**, 305 CL terms collapsed to 21 coarse groups, **17 seconds of CPU** to train, 38 MB. Two mouse datasets were excluded from the reference entirely and used as the test set — 12,646 cells, 137 truth types, 41 tissues, none of it seen in training:

	ontology	coverage	cells/s
all cells	0.638	100%	9,712
confidence ≥ 0.5	0.718	71%	

Table 8. `broad_mouse_v1` on 12,646 held-out mouse cells (137 truth types, 41 tissues, none of it in the reference), with and without abstain.

Its abstain calibration is also better behaved than the human census model’s — 0.88 accuracy at 82% coverage, against 0.80 at 46% — because mouse census carries fewer near-duplicate subtypes than human’s ~800-way vocabulary. Read across organisms with care: these are different queries, so the mouse numbers sit in the same regime as the human ones rather than being comparable to them. Two limits apply. The ablation above was run on **human** and applied to mouse; CL is species-neutral by construction, but nothing here shows it groups mouse types as well. And mouse census is shallow in *datasets* — 51 in total, one embryo atlas holding 11.4M of 18.4M cells — so tissue breadth is good while lab and protocol breadth is far below human’s 487 datasets.

The focused pass — tissue-specific. For the tissue the broad pass identifies, a small focused reference re-annotates at full resolution. The gain is large and it is the crux of the workflow: on held-out **cross-study** liver cells, the broad census model scores exact-CL **0.23** / ontology **0.58**, while a focused **48-type HLiCA liver** reference on the *same cells* reaches **0.72** / **0.86**. Refinement is where fine-grained accuracy comes from; the broad model’s job is to route to it, not to be right about subtypes itself.

The broad pass hands the query to the focused pass; the two do not combine. The natural assumption is that a better broad call should also make the focused call better. It does not. On the leakage-free cross-study liver split, substituting the stronger **Pan-human Azimuth** for the broad pass lifts it (ontology 0.380 vs 0.338 for our own broad model) but changes nothing downstream. Using that broad call to *narrow* the focused pass’s classes — the zero-retrain masking `actinn-jax` ships — makes the result **worse**, $0.731 \rightarrow 0.708$: the broad call matches the true lineage on 85.8% of cells, and the 14% it misses cost more than the 86% it gets right can gain, since a wrong mask discards the correct class outright. A **perfect** coarse call would add only **+2.8 points** (0.759). Once the focused pass covers the tissue, there is almost no accuracy left for a broad model to contribute; its value is choosing which focused reference to load and covering cells none of them claims (`PAN_HUMAN_AZIMUTH.md`).

Supporting mechanisms. (i) A **foundation-model-guided coarse→fine hierarchy**: a foundation model discovers a coarse→fine structure *offline*, the fast CPU model uses it at inference — beats a flat classifier on 3/4 datasets, matches an expert hierarchy, beats a random-grouping control, and never depends on the foundation model at prediction time. (ii) **Within-cell-type resolution**: the same machinery resolves hepatocyte zonation (portal→central) at ~0.99 within-one-zone, generalizing across donors and datasets — a third stage below the cell-type label. Together these make `actinn-jax` a *pipeline* (broad $\rightarrow$ tissue $\rightarrow$ subtype/state), not just a classifier, at classical-method speed throughout.

3.5 Rejection / abstain

Holding out 9 of 36 cell types entirely from the HLiCA liver reference (so 1,350 query cells are genuinely out-of-distribution), and sweeping a confidence threshold:

threshold	actinn-jax: acc(kept) / coverage / OOD-flagged	CellTypist: acc(kept) / coverage / OOD-flagged
0.0	0.885 / 1.00 / 0.00	0.873 / 1.00 / 0.00
0.5	0.919 / 0.93 / 0.30	0.936 / 0.68 / 0.68
0.7	0.946 / 0.84 / 0.52	0.936 / 0.68 / 0.68
0.9	0.969 / 0.66 / 0.73	0.937 / 0.68 / 0.68

Table 9. Confidence-threshold sweep with 9 of 36 cell types held out of the liver reference, so 1,350 query cells are genuinely out-of-distribution. Each cell reads accuracy on kept cells / fraction of cells kept / fraction of held-out-type cells flagged.

actinn-jax’s probability gives a **smooth, tunable** abstain curve — accuracy on kept cells rises 0.885→0.969 as coverage falls 1.00→0.66 and OOD-flagging climbs 0.00→0.73 — so a user can dial the precision/coverage/novel-cell-detection trade-off. CellTypist’s probabilities are **saturated** (near 0 or 1): the threshold jumps to 68% OOD-flagged at 0.5 and then does not move across 0.5–0.9, giving essentially one operating point rather than a tunable curve. scmap-cluster offers a single native “unassigned” decision (no sweep). A tunable threshold is what makes the abstain step usable as a routing decision in the workflow of §3.4.

3.6 Foundation-model zero-shot (scPRINT)

scPRINT is run separately as a zero-shot predictor (no training on the reference), as a reference point for the “just use a foundation model” alternative. On the lung dataset it scored **exact accuracy 0.029 / ontology concordance 0.201, taking 62 s for 2,694 query cells** — against actinn-jax’s 0.894 in 0.23 s — so as a *label* predictor it is both slow (~280× actinn-jax) and weak. It is reported on lung alone because its fixed vocabulary cannot score the other datasets: the liver and blood+gut sets contain Cell Ontology terms absent from its label set, and it declines them rather than guessing. Its pretrained classifier head also carries a **fixed label vocabulary** and requires CL ids, so it cannot emit labels outside that set; this is an inherent property of a zero-shot label head, not a defect, but it means the foundation model’s raw predictions are not a drop-in annotator. This is exactly why our two-stage workflow uses scPRINT’s **embeddings** — its learned structure — to shape a small trained model, never its zero-shot labels: the foundation model is valuable, its raw label predictions are not.

3.7 External validation: Open Problems label_projection

Our in-house benchmark is neutral but self-run. For an **independent** check, we ran actinn-jax through the community-standard [Open Problems](#) label-projection task — their 6 datasets, splits, and metrics — and slotted it into their published v2.0.0 leaderboard.

Placement. On a benchmark we did not design and cannot tune, **actinn-jax places 3rd of 17 on mean accuracy (0.837), 1st among all methods that complete every dataset**, beats its PCA-space mlp sibling, and posts the best accuracy of any completing method on the hardest dataset (tabula_sapiens, 160 types: 0.394 vs mlp 0.342). It is upper-mid on macro-F1 and **not the top method overall** — scANVI+scArches surgery and xgboost [Chen & Guestrin 2016](#) score higher on the datasets they finish, though both **fail to complete tabula_sapiens**, which is why they rank above actinn-jax only on a mean over the 5 easier datasets. The foundation models again

land at the bottom (scgpt_zeroshot 0.639, uce 0.131 $\approx$ the random-labels control), reproducing §3.6 externally.

Controlled same-hardware speed & memory. Leaderboard runtimes mix hardware, so we re-ran actinn-jax and the CPU tier through OP’s *own* Nextflow pipeline on a **single AWS r7i.8xlarge** — identical hardware, harness, storage, and trace instrumentation for every method. This is the paper’s cleanest cross-method cost comparison:

method (same box)	mean acc	macro-F1	runtime/dataset	peak RSS
mlp	0.838	0.664	327 s	19.9 GB
actinn-jax	0.837	0.655	165 s	21.0 GB
logistic_regression	0.813	0.689	29 s	20.0 GB
xgboost	0.795	0.613	905 s	80.7 GB
knn	0.793	0.648	11 s	19.5 GB
cellmapper_linear	0.776	0.553	58 s	31.5 GB
naive_bayes	0.738	0.613	31 s	19.5 GB

Table 10. Open Problems methods re-run on a single AWS r7i.8xlarge — identical hardware, harness, storage, and trace instrumentation for every method — averaged over the benchmark’s datasets, ordered by accuracy as in Table 3.

actinn-jax **ties its mlp sibling on accuracy at $\sim 2\times$ the speed** (165 vs 327 s) and **beats xgboost’s accuracy at $\sim 5.5\times$ less runtime and $\sim 4\times$ less memory** (165 s / 21 GB vs 905 s / 81 GB). The faster methods (knn, logistic) are less accurate; the more accurate heavyweight (xgboost) is far slower and heavier: **within Open Problems’ method set**, on controlled hardware, nothing beats actinn-jax on accuracy and cost at the same time. That set contains no carefully-tuned linear pipeline; we added one and ran it on the same six datasets, and it does not change the ranking (Table 11).

Extension: our four added methods on the same datasets. We ported the tuned linear pipeline, the SGD-SVM, CellTypist and scTOP into Open Problems components and ran all four on all six datasets (per-run numbers in the [benchmark repository](#)). Accuracy and macro-F1 do not depend on the machine, so these were run locally rather than on the rented instance; **cost is not reported for them**, because that does depend on the machine and only Table 10 was measured under controlled conditions. The premise is checkable rather than assumed: running actinn-jax’s component locally reproduces its controlled-run accuracy on every dataset to within **0.004** (mean difference 0.0005; three of six identical to four decimals), which is smaller than the spread between repeats.

method	mean acc	macro-F1	dkd	gtex	hypomap	immune	pancreas	tabula
linear-anova-pca	0.829	0.647	0.950	0.870	0.996	0.834	0.968	0.356
SVM (SGD)	0.815	0.651	0.939	0.838	0.996	0.839	0.965	0.310
CellTypist	0.809	0.644	0.938	0.837	0.996	0.843	0.964	0.275
scTOP	0.581	0.462	0.898	0.124	0.991	0.495	0.939	0.042

Table 11. The four methods we added to Open Problems, on all six of its datasets, ordered by accuracy. Dataset columns abbreviate the names used in Table 10 (immune = immune_cell_atlas, pancreas = mouse_pancreas_atlas, tabula = tabula_sapiens). Comparable to Table 10 on accuracy, which is machine-independent; not on cost, which is not measured here.

Read against Table 10, **the tuned linear pipeline does not repeat its §3.1 win here:** 0.829 places it third, behind `mlp` (0.838) and `actinn-jax` (0.837) and ahead of OP’s own `logistic_regression` (0.813). The SGD-SVM (0.815) and CellTypist (0.809) land in the same band. The linear pipeline’s lead on our own panel (§3.1) therefore does not generalize: whichever of the two panels is asked, the answer is a cluster of methods within a couple of points of each other, but which one tops the cluster changes with the datasets.

scTOP fails on two of the six. Its mean (0.581) is not a uniformly weak result but two collapses inside four ordinary ones: 0.124 on `gtex_v9` and 0.042 on `tabula_sapiens` against 0.90–0.99 elsewhere. Open Problems hands every method 1,000 HVGs, and scTOP’s rank projection needs genes expressed in an appreciable fraction of cells; after its expression filter only **106 of 1,000** genes survive on `gtex_v9` and 210 on `tabula_sapiens`, against 306–405 where it works. At that point some test cells retain no counts at all — 185 of 11,508 on `gtex_v9` — and cannot be placed; they are scored as wrong rather than dropped, so 0.124 is a floor. This is the same limitation §3.1 reports from the other direction: scTOP is built for small, low-cardinality problems, and a fixed 1,000-HVG budget over 53 and 160 types is neither.

Three cheap ablations, and their limits. (i) *Input standardization* — z-scoring genes to the reference’s frozen mean/std, a CPU-only domain-alignment inspired by scArches reference surgery — lifts mean accuracy +0.2 and macro-F1 +1.2 pt (largest on batch-shifted datasets: `gtex` macro-F1 +7.9). It is shipped **off by default** because it shifts the softmax calibration that the two-stage abstain thresholds are tuned against. (ii) *Gene budget* — OP feeds every method 1000 HVGs, which most starves a gene-space MLP. Widening to ~5000 HVGs lifts accuracy on 4 of 6 datasets (immune/`gtex` +3 pt) and there **matches scANVI+scArches’s full-config leaderboard accuracy** (immune 0.891 $\approx$ 0.892) on CPU in minutes — the gap to the top method was largely a gene-budget artifact, not the VAE. But more genes is **not** a universal win: it *regresses* `tabula_sapiens` by ~10 pt (its 284-cell test batch across 160 fine types overfits reference-specific genes) and saturates hypomap (Figure 4). The budget is **selectable without test labels**: held-out *reference* cross-validation rises for the datasets that benefit and is the one signal that *drops* for `tabula_sapiens`, and a trivial query-cells-per-class check independently flags it (Figure 5) — so the budget can be set deterministically per dataset. (iii) *Negative control* — a CPU-only, UCE-style protein-embedding featurization (expression-weighted mean of ESM2 gene embeddings) does **not** help and hurts the hardest case: the pooling discards the per-gene detail that separates fine types, so the value of a foundation model like UCE stays locked in its GPU transformer, not a portable averaging trick.




Figure 4. actinn-jax accuracy and macro-F1 against input gene budget across all six Open Problems datasets. More genes help most datasets but regress the fine-grained, domain-shifted `tabula_sapiens`.




Figure 5. Label-free signals for setting the gene budget without test labels. Held-out reference cross-validation and query-cells-per-class both single out `tabula_sapiens`, the one dataset where more genes hurt.

4 Discussion

Method choice is now dominated by cost, not accuracy. The five leading methods are separated by 0.008 in mean accuracy and by $\sim 130\times$ in inference time (§3.1), so for most annotation jobs the decision is made on the cost axes. On laptop-sized references the tuned linear pipeline is the best of them on accuracy-per-second (0.839, fitting in 4.0 s); at atlas scale ProtoCloud is the most accurate by a clear margin (0.976 vs 0.936 at 49k lung cells; 0.905 vs 0.824 at 47k liver cells), having been below the cluster on the subsampled matrix. `actinn-jax`’s place in that picture rests on a cost profile with two properties that matter for repeated use rather than for a single labelling run.

Footprint. Inference is **sub-second and flat** — independent of reference size and cardinality (§3.3) — so a query costs the same against a 1k-cell reference and a 49k-cell one. Peak memory is unremarkable on small references (2.1 GB mean — seven of the ten other methods are lighter), but it is the lowest of those we carried to **atlas scale**: $\sim 2\times$ below the linear pipeline (6.1 vs 13.2 GB at 49k lung cells; 6.5 vs 12.6 GB at 47k liver cells), and below scTOP above ~ 25 k cells, where scTOP’s rank processing densifies and crosses over (9.3 vs 6.5 GB at 47k). The advantage is bounded at $\sim 2\text{--}3\times$ rather than widening. Combined with the train-once/map-many cache, a reused reference pays fit once and then only the flat predict cost, where the linear pipeline refits scaler/PCA/classifier for every query.

Those two properties are what the workflows are built on. The practical gains come not from the flat classifier but from the **large**→**refined** pipeline (§3.4): a census-scale broad model with a calibrated abstain routes a query to a small focused reference that re-annotates at full resolution (cross-study liver 0.23/0.58 $\rightarrow$ 0.72/0.86), with zonation as a further stage. The pipeline is only usable because each hop is cheap — a flat sub-second call against a cached model, at bounded memory — so several models can be chained, and a large reference kept resident, on a laptop without a GPU. A method that is a few points more accurate but pays a full fit per query, or a multi-gigabyte dense expansion per stage, does not compose the same way.

Where it loses, and why. actinn-jax trails on small references and very-few-cells-per-type sets (§3.1, §5), and — a finding from the OP analysis — it is sensitive to input budget under domain shift: more genes help most datasets but overfit a tiny, fine-grained, shifted query (tabula_sapiens). The two cheapest levers are principled and label-free: **standardization** (a scArches-style domain alignment) and a **gene budget chosen by cross-validation on the reference alone** both improve the model without peeking at test labels, and at a fair gene budget the gene-space MLP *matches the leaderboard’s top method* on the datasets both complete — In other words, much of what looked like a deficit of the *model* was a deficit of its **input**: the benchmark hands every method 1,000 highly variable genes, which is a far harsher constraint on a model reading gene space directly than on one that learns a compressed representation first. Widen the gene panel and most of the gap closes — so the quantity being measured was partly “how well does each method tolerate a narrow input”, not “which model class is stronger”.

Foundation models. Their **zero-shot labels** are the weakest option in both benchmarks (scPRINT, scGPT, UCE $\approx$ random); their **embeddings/structure** are where the value lies, and our two-stage hierarchy uses exactly that. A cheap CPU shortcut to a foundation-model representation (protein-embedding pooling) did *not* transfer (§3.7), underscoring that the value is in the trained transformer, not a portable trick. The concurrent **Pan-human Azimuth** effort reaches the same conclusion from a far larger curated corpus: a 7M-parameter supervised hierarchical MLP over a 5,055-gene panel yields cleaner annotations than scGPT and SCimilarity, whose labels they find fragmented and mis-transferred. Their stated theme — that training-data quality and organization matter as much as architecture or scale — and their finding that accuracy **saturates past ~ 5 M training cells** (PAN_HUMAN_AZIMUTH.md) are independent support, from a group with no stake in our conclusion, for both halves of our reading.

When extra expressivity pays. Two lines of evidence bear on this. First, **ProtoCloud** — a prototype-based, self-explaining VAE with built-in uncertainty and gene-level attribution, a strictly richer model than ours — does not sit a fixed distance above or below us. Where it lands depends on how much reference data it is given: below the top cluster on the subsampled matrix, ahead on lung, and the strongest method of all at atlas scale (§3.1, §5), at $\sim 9\times$ actinn-jax’s CPU fit time (146 vs 16 s mean) (PROTOCLOUD.md). The richer model pays off where it has the data to support

it, and not before — which is an argument for matching model capacity to reference size, not for preferring small models everywhere. Second, and more broadly, Souza & Mehta [Souza & Mehta 2026](#) report that **parameter-free linear representations** match or exceed single-cell foundation models across cross-species transfer, human cell-type classification and disease-state prediction — on Tabula Sapiens 2.0 their normalize→ANOVA→PCA→logistic-regression pipeline reaches mean macro-F1 0.899 against 0.907–0.910 for TranscriptFormer variants — and they offer a geometric explanation: the biologically realized cell manifold is well approximated by a **linear subspace**, so once noise is suppressed performance *saturates* and extra expressivity buys little. That is a principled account of why a four-layer MLP over normalized gene space stays competitive, and it matches our own finding that the residual gap to heavier methods was mostly **the size of the input each method was given rather than the class of model** (§3.7). They further find that simple methods hold up *best* out-of-distribution — on novel cell types and unseen organisms. That matters for the abstain mechanism of §3.5, which is only useful if low confidence tracks genuine novelty rather than model noise. A method that degrades gradually on unfamiliar data still assigns lower confidence to the cells it is getting wrong, so a threshold separates them; a method that degrades abruptly can be confidently wrong, and no threshold helps.

Practical guidance. For a **one-off annotation** on laptop-sized data, use the **tuned linear pipeline** (normalize → ANOVA → PCA → logistic regression): it is the most accurate-per-second method here. With a **real atlas and time to train**, **ProtoCloud** is the most accurate method we benchmarked (0.976 at 49k cells), and the GPU it targets makes its 19× CPU fit cost much less punishing in practice. **scTOP** suits small, low-cardinality problems, where it is by far the cheapest (~1 s) and competitive. **actinn-jax** fits the cases its cost profile is shaped for: when **memory is the binding constraint** (~2× lighter than the linear pipeline at scale, with memory-bounded chunked inference), when a reference is **reused** across many queries (the cached model amortizes the fit), and when the job is the **multi-stage workflow** rather than a single label — broad→refined hand-off, tissue-aware refinement, calibrated abstain, novel-cell-type screening.

For the **broad pass specifically**, prefer a **purpose-built pan-human model**: **Pan-human Azimuth** (PAN_HUMAN_AZIMUTH.md) ships a pretrained hierarchical classifier over a harmonized organism-wide typology (8 levels, 382 leaf types, trained on 9.7M curated cells), with abstention learned rather than thresholded and an **expected calibration error (ECE)** of 0.0044 — meaning its stated confidence sits within half a percentage point of its observed accuracy — running at ~1,000 cells/s on a laptop. It is better resourced than our census-built reference and we make no claim to improve on its annotations. What we do with it instead is **use it**: distilling its labels and its hierarchy into actinn-jax produces a broad-pass model of comparable accuracy at roughly nine times its throughput, built without a GPU and without labeled data (§3.4). A curated pan-human model is the right thing to start from; a small fast model is the right thing to iterate with.

The part of any fixed typology that cannot be performed is the **hand-off**: it cannot re-annotate into a user’s own focused label set (the HLiCA liver reference, cross-study 0.23/0.58 → 0.72/0.86) or resolve states below a leaf (hepatocyte zonation). ProtoCloud likewise provides uncertainty, gene attribution and a refinement path (fine-tuning a pretrained model). What survives as distinct to actinn-jax is therefore not the breadth of the broad reference but **the chain it enables** — start pan-human, re-annotate into label sets no pretrained model carries, screen what none of them claims — at a cost that keeps every stage on a laptop. The distillation recipe is not specific to this teacher or to humans: it needs a pretrained annotator whose typology can be read off its outputs, so extending the same entry point to other organisms is a matter of finding one (or building the broad-pass reference from a labeled census slice, as §3.4’s census route does).

What the speed is for. Annotation is almost never the result; it is the step before the biology, and it is usually run more than once — the labels change when a reference is swapped, a threshold moves, a new sample arrives, or a cluster turns out to be two things. The argument of this paper is not that a faster classifier produces better labels; at equal data most of these methods produce much the same labels (§3.1). It is that when a stage costs a second instead of a minute, and a chain of stages fits in memory on the machine already on the desk, the loop from question to annotated data to the next question closes in an afternoon rather than across a queue. That is the contribution we would defend: not a better number in a table, but a workflow — start from the best pan-human model available, refine into the labels a given study actually uses, and flag what neither covers — that a working scientist can run, inspect, and run again.

5 Limitations

- Single hardware family for the in-house panel (Apple Silicon); no discrete-GPU numbers there. The deep and foundation methods are written for CUDA and we would expect them to be faster on it, but we did not measure that — it is outside the “runs-on-a-laptop” question, and every timing here should be read as this hardware’s. The §3.7 controlled run covers the **CPU tier** on one cloud box; the GPU/R methods there are reported from OP’s own cloud-CI trace (indicative). A same-hardware GPU-tier run to fold those into a single controlled table is the natural extension.
- **Wall-clock comparisons across methods are only meaningful at matched concurrency.** Measured %cpu on the OP harness spans 49% to 2338% across methods — some are single-threaded, some saturate 23 cores — so the same method timed at two concurrency settings differs by up to 12×, and any wall-clock ranking taken at high concurrency silently ranks threading and scheduler contention alongside algorithm. Cost is therefore reported only from the single matched-concurrency run (Table 10), which covers OP’s own seven methods; the four we added are compared on accuracy alone (Table 11). Putting all eleven on one cost axis needs one more controlled run.
- Subsampled references (to keep the matrix tractable); the scaling section characterizes the size dependence directly.
- **The cross-method comparison is human only;** six datasets per benchmark; GPU foundation models beyond scPRINT/UCE (scGPT, Geneformer, popV) not run locally. The shipped references now cover mouse (§3.4), but no *method comparison* was run on mouse data — the pan-mouse result establishes that the reference-building route works on a second organism, not that actinn-jax’s standing against other methods carries over to it.
- actinn-jax needs more cells/type than linear methods to reach parity (§3.1); on very small references it trails.
- **The distilled broad-pass reference inherits a vocabulary, not a calibration.** Pan-human Azimuth’s abstention is trained (and its ECE measured at 0.0044); the distilled student learns hard labels only, so it carries none of that — its confidence barely separates right from wrong (§3.4). Recalibrating it, or distilling from soft targets, is the obvious next step and would need a loss change in the package. Its evaluation is also two withheld *atlases*, not withheld tissue: the census-wide corpus spans 376 tissues, so nothing here measures behaviour on biology the corpus never saw.
- **The gene budget is dataset-dependent, not a free parameter.** More genes help most references but *overfit* a tiny, fine-grained, domain-shifted query (tabula_sapiens −10 pt); our reference-cross-validation + cells-per-class selection rule is validated on 6 datasets with a

single clean failure case, so it is directional evidence, not a tuned threshold.

- **Standardization is shipped opt-in**, not default, because it shifts probability calibration that the two-stage abstain thresholds are tuned against; combining the two cleanly would require re-tuning those thresholds.
- **Our classical tier (SVM, kNN, CellTypist) is untuned, and a tuned linear baseline beats actinn-jax.** Souza & Mehta’s pipeline (normalize → ANOVA → standardize → PCA(220) → logistic regression) leads the matrix on accuracy (0.839 vs 0.831) and macro-F1 (0.699 vs 0.683) while fitting 7× faster on this panel’s hardware (§2.5), with its largest margin on the 86-type blood+gut set (+4.2 pt) (`SIMPLE_BASELINES.md`). Wall-clock ratios between a JAX model and a scikit-learn pipeline shift with the CPU and the threading available, so read them as this machine’s, not as constants. The §3.1 margins over the classical tier should therefore be read as margins over *untuned* baselines; a like-for-like tuning effort on SVM or CellTypist would likely narrow them further. scTOP is cheaper still to fit but competitive only at small, low cardinality — best pbmc macro-F1 (0.837), degrading to 0.649 on liver.
- **actinn-jax’s memory advantage is bounded, and the linear pipeline scales further than a naive extrapolation implies** (`SCALING_MEMORY.md`). Extrapolating the linear pipeline’s small-reference footprint overstates its atlas-scale peak by ~2.4× — the measured figure is **13.2 GB** at 49k reference cells — because its feature-selection step caps the dense matrix at 20,000 genes rather than the full gene set. The linear/actinn-jax memory ratio is likewise **bounded at ~2–3×** (2.15× at full scale) rather than widening with data, so there is no scaling cliff and no regime tested where the linear pipeline stops being usable. Memory in this setting should be measured, not extrapolated.
- **At atlas scale the ordering changes and the deep model leads.** Scaling the lung reference 3k → 49k cells, **ProtoCloud goes from worst (0.722) to best (0.976)**, clear of actinn-jax (0.936) and the linear pipeline (0.939), at 19× the CPU fit cost; scTOP gains nothing from scale (0.819 → 0.846) and loses its memory edge (1.22× actinn-jax at 49k). This replicates on the HLiCA liver atlas (36 types, 525k cells, sweep to 47k ref): same bounded ~2× memory band, same ProtoCloud reversal (0.581 → 0.907), same scTOP stagnation (0.657 → 0.583), with scTOP crossing over to heavier than actinn-jax (1.41×). Conclusions drawn from subsampled references do not transfer to atlas scale, in either direction.
- **Annotation only.** Cross-species transfer and disease-state prediction — the tasks where Souza & Mehta find the largest foundation-model deficits, and where a fast method would be most attractive — are outside this benchmark’s scope (human, within/cross-dataset annotation).

Data & code availability

Two repositories hold everything needed to reproduce this work. The [benchmark repository](#) carries the harness, the per-method adapters, the run configurations that produce every number reported here, the environment lockfiles, the result tables including the unified eleven-method matrix, and the figure scripts. The [actinn-jax package](#) is on PyPI (`pip install actinn-jax`). Rebuilding the shipped broad reference is a single documented command; the distilled broad-pass reference, the pan-mouse reference and the Cell-Ontology hierarchy it depends on each have their own build documentation in the benchmark repository, and the file-level detail lives there rather than here.

Pre-trained references — human (`broad_human_v1`, `panhuman_distill_v1`), mouse (`broad_mouse_v1`) and focused liver (`liver_hliCA_v1/v2`) — are archived at [doi:10.5281/zenodo.10000000](https://doi.org/10.5281/zenodo.10000000)

[zenodo.21688151](#) (CC BY 4.0; cite the concept DOI [10.5281/zenodo.21688150](#) for the latest version). The package downloads and caches them on first use, so no manual retrieval is needed, and can pre-download them for offline use.

HLiCA data © Edgar et al. 2026 (CC-BY 4.0), [doi:10.64898/2026.06.30.735539]. The distilled reference derives from **Pan-human Azimuth** (Sarkar, Li, Molla, ... Satija, bioRxiv 2026, doi:10.64898/2026.07.16.738997); its weights are © the authors under CC BY 4.0, obtained via **panhumanpy** (MIT) and [Zenodo](#). Attribution is a condition of that licence and travels inside each shipped model's build record.

References

- 10x Genomics. 3k PBMCs from a healthy donor, Cell Ranger 1.1.0 (2016). Distributed with scanpy as pbmc3k. [10xgenomics.com/datasets/3-k-pbm-cs-from-a-healthy-donor-1-standard-1-1-0](#).
- Abdelaal T, et al. A comparison of automatic cell identification methods for single-cell RNA sequencing data. *Genome Biology* 20:194 (2019). doi:10.1186/s13059-019-1795-z.
- Alegbe T, Harris BT, Fachal L, et al. Cell-type-resolved genetic variation shapes inflammatory bowel disease risk (IBDverse). *Nature* (2026). doi:10.1038/s41586-026-10627-z.
- Aran D, et al. Reference-based analysis of lung single-cell sequencing reveals a transitional profibrotic macrophage (SingleR). *Nature Immunology* 20:163-172 (2019). doi:10.1038/s41590-018-0276-y.
- Bradbury J, et al. JAX: composable transformations of Python+NumPy programs (2018). [github.com/jax-ml/jax](#).
- CZI Cell Science Program. CZ CELLxGENE Discover Census, LTS release 2025-11-08. [chanzuckerberg.github.io/cellxgene-census](#).
- Chen T, Guestrin C. XGBoost: a scalable tree boosting system. *KDD* 785-794 (2016). doi:10.1145/2939672.2939785.
- Domínguez Conde C, et al. Cross-tissue immune cell analysis reveals tissue-specific features in humans (CellTypist). *Science* 376:eabl5197 (2022). doi:10.1126/science.abl5197.
- Edgar RD, Portman JR, Hu H, et al. HLiCA: an integrated cell atlas of the healthy human liver. *bioRxiv* (2026). doi:10.64898/2026.06.30.735539.
- Fu Q, Dong C, Liu Y, et al. A comparison of scRNA-seq annotation methods based on experimentally labeled immune cell subtype dataset. *Briefings in Bioinformatics* 25(5):bbae392 (2024). doi:10.1093/bib/bbae392.
- Guo K, Ding J. ProtoCloud: a prototypical self-explaining model for single-cell analysis. *Cell Genomics* 6(6):101217 (2026). doi:10.1016/j.xgen.2026.101217.
- Kalfon J, Samaran J, Peyré G, Cantini L. scPRINT: pre-training on 50 million cells allows robust gene network predictions. *Nature Communications* 16:3607 (2025). doi:10.1038/s41467-025-58699-1.
- Kiselev VY, Yiu A, Hemberg M. scmap: projection of single-cell RNA-seq data across data sets. *Nature Methods* 15:359-362 (2018). doi:10.1038/nmeth.4644.
- Lin Z, et al. Evolutionary-scale prediction of atomic-level protein structure with a language model (ESM-2). *Science* 379:1123-1130 (2023). doi:10.1126/science.ade2574.
- Lotfollahi M, et al. Mapping single-cell data to reference atlases by transfer learning (scArches). *Nature Biotechnology* 40:121-130 (2022). doi:10.1038/s41587-021-01001-7.
- Ma F, Pellegrini M. ACTINN: automated identification of cell types in single cell RNA sequencing. *Bioinformatics* 36(2):533-538 (2020). doi:10.1093/bioinformatics/btz592.

17. Open Problems for Single-Cell Analysis Consortium. Open Problems: a living benchmark for single-cell analysis (2024). openproblems.bio.
18. Pedregosa F, et al. Scikit-learn: machine learning in Python. *JMLR* 12:2825-2830 (2011). jmlr.org/papers/v12/pedregosa11a.html.
19. Rosen Y, et al. Universal cell embedding provides a foundation model for cell biology (UCE). *Nature* (2026). [doi:10.1038/s41586-026-10689-z](https://doi.org/10.1038/s41586-026-10689-z).
20. Sarkar S, Li Z, Molla G, et al. Organism-scale annotation with Pan-human Azimuth. *bioRxiv* (2026). [doi:10.64898/2026.07.16.738997](https://doi.org/10.64898/2026.07.16.738997).
21. Sikkema L, et al. An integrated cell atlas of the lung in health and disease (HLCA). *Nature Medicine* 29:1563-1577 (2023). [doi:10.1038/s41591-023-02327-2](https://doi.org/10.1038/s41591-023-02327-2).
22. Souza H, Mehta P. Parameter-free representations outperform single-cell foundation models on downstream benchmarks. *bioRxiv* (2026). [doi:10.64898/2026.02.11.705358](https://doi.org/10.64898/2026.02.11.705358).
23. Travaglini KJ, Nabhan AN, Penland L, et al. A molecular cell atlas of the human lung from single-cell RNA sequencing. *Nature* 587(7835):619-625 (2020). [doi:10.1038/s41586-020-2922-4](https://doi.org/10.1038/s41586-020-2922-4).
24. Xu C, et al. Probabilistic harmonization and annotation of single-cell transcriptomics data with deep generative models (scANVI). *Molecular Systems Biology* 17:e9620 (2021). [doi:10.15252/msb.20209620](https://doi.org/10.15252/msb.20209620).
25. Yampolskaya M, Herriges MJ, Ikonomou L, Kotton DN, Mehta P. scTOP: physics-inspired order parameters for cellular identification and visualization. *Development* 150(21):dev201873 (2023). [doi:10.1242/dev.201873](https://doi.org/10.1242/dev.201873).
26. Munroe R. Standards. *xkcd* 927. xkcd.com/927.